

BERT & its family

Lecturer: Hiếu Trần

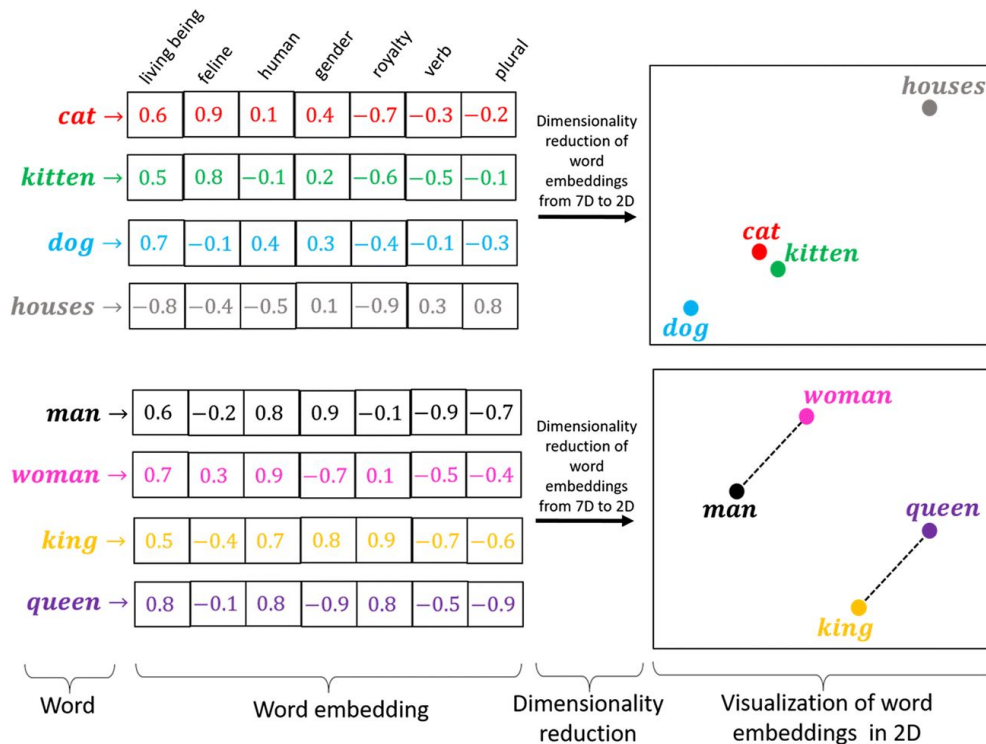
OUTLINE

- What is Contextualized Word Embedding?
- Types of Learning
- Pre-trained Language Models
- Downstream Tasks

What is Contextualized Word Embedding?

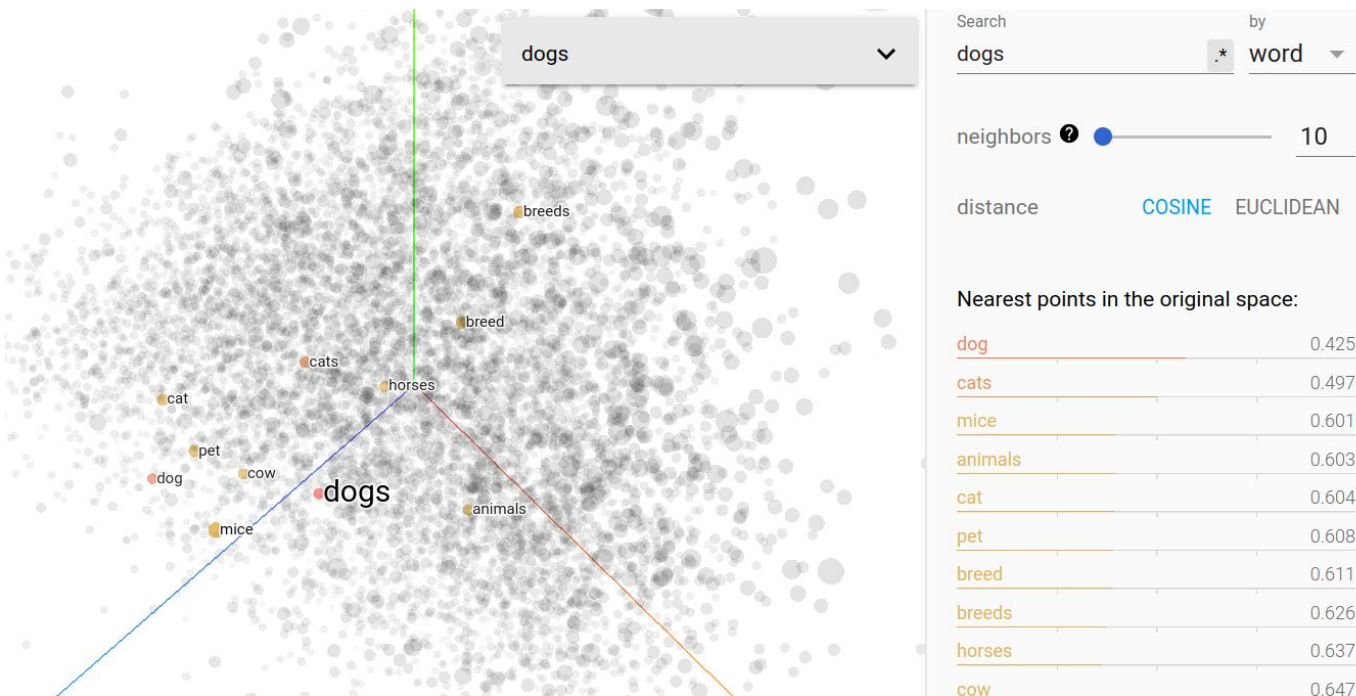
WORD EMBEDDING

- Word embeddings are the basis of deep learning for NLP
- Word embeddings (**word2vec**, **GloVe**) are often *pre-trained* on text corpus from co-occurrence statistics



WORD EMBEDDING: Examples

- Visualizing Word2Vec using Tensorflow: <https://projector.tensorflow.org/>



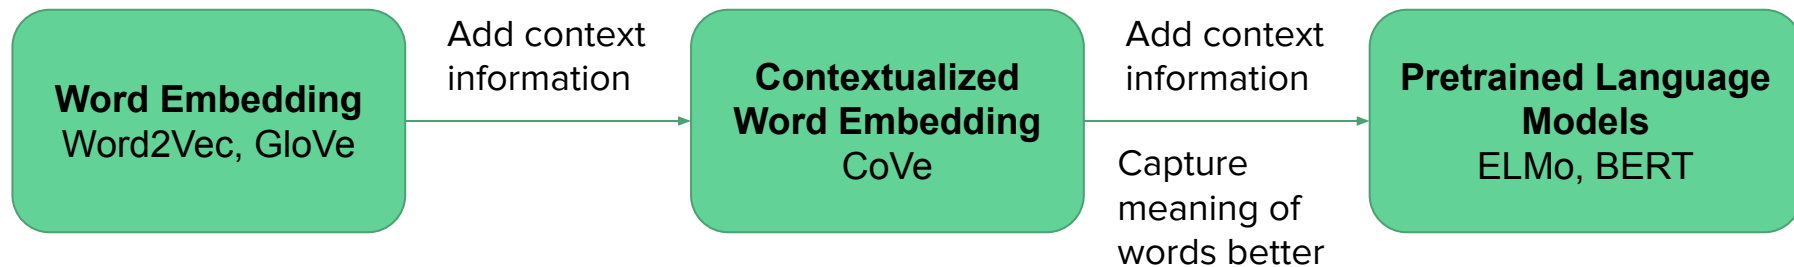
WORD EMBEDDING: Limitations

- TF-IDF: based on the frequency of words and the rarity of the words, not leverage the context i.e. word co-occurrence
- Word2Vec: based on word co-occurrences in local context
- GloVe: based on word co-occurrences in global context

Limitation: Don't have context information!!!

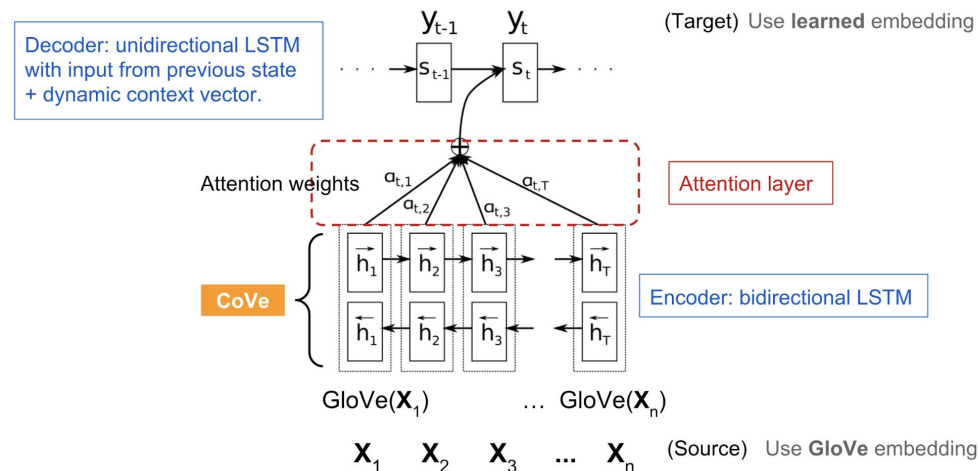
- Combine all the different meanings of the word into one vector
 - Anh ta sử dụng **kiểm** rất điều luyện.
 - **Kiểm** ăn bây giờ khó lắm.

IMPROVEMENTS IN NLP




CONTEXTUAL WORD VECTORS (CoVe)

- CoVe:
 - Learn word embedding by an encoder in seq-to-seq (biLSTM) machine translation model
 - Use GloVe as an initial word embedding
- Need supervised data



EMBEDDINGS FROM LANGUAGE MODEL (ELMo)

- ELMo:
 - Predict a token based on the history
 - Concatenation of right-to-left and left-to-right LSTMs




Types of Learning

TYPES OF LEARNING


- **Multi-task learning** is a general term for training on multiple tasks.
- **Transfer learning** is a type of multi-task learning where we only really care about one of the tasks.
- **Pre-training** is a type of transfer learning where one objective is used first.

WHAT IS MULTI-TASK LEARNING?




- Train representations to do well on multiple tasks at once.
- When to use multi-task?
 - When one of your tasks has many fewer data.
e.g., high-resource language → low-resource language
 - When your tasks are related.
e.g., predicting POS tagging and punctuation.

WHAT IS MULTI-TASK LEARNING?



- Train representations to do well on multiple tasks at once.
- When to use multi-task?
 - When one of your tasks has many fewer data.
e.g., high-resource language → low-resource language
 - When your tasks are related.
e.g., predicting POS tagging and punctuation.


WHEN USE MULTI-TASK LEARNING?



- When one of your tasks has many fewer data.
e.g., high-resource language → low-resource language
- When your tasks are related.
e.g., predicting POS tagging and punctuation.

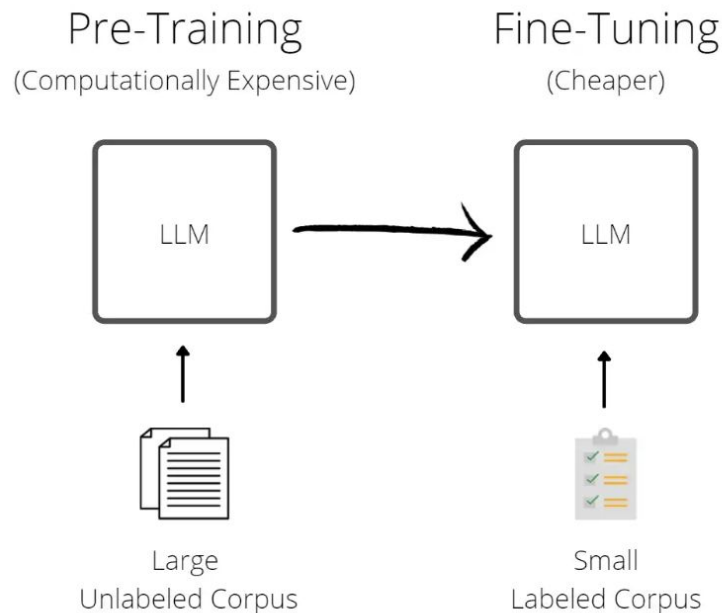
WHAT IS PRE-TRAINING?

- First train on one task, then train on another.
- Widely used in word embeddings, contextualized word representations.



WHY NEED PRE-TRAINING?


- Pre-training: train on **unlabeled data** over different pre-training objectives
 - Self-supervised pre-training
- Fine-tuning: tune parameters using **labeled data** from the downstream tasks



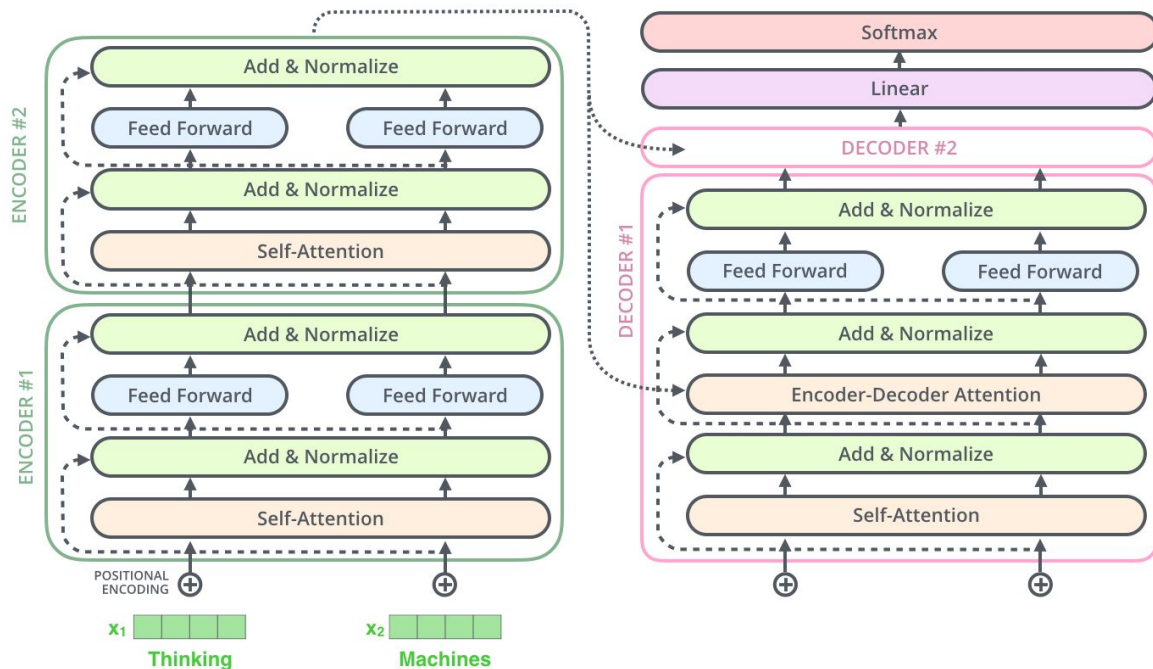
WHAT MAKE PRE-TRAINED MODELS?

Pre-training methods often refer to a combination of

- **Model:** The underlying neural network architecture
- **Training Objective:** What objective is used to pretrain
- **Data:** What data the authors chose to use to train the model



Transformer Architecture: Review

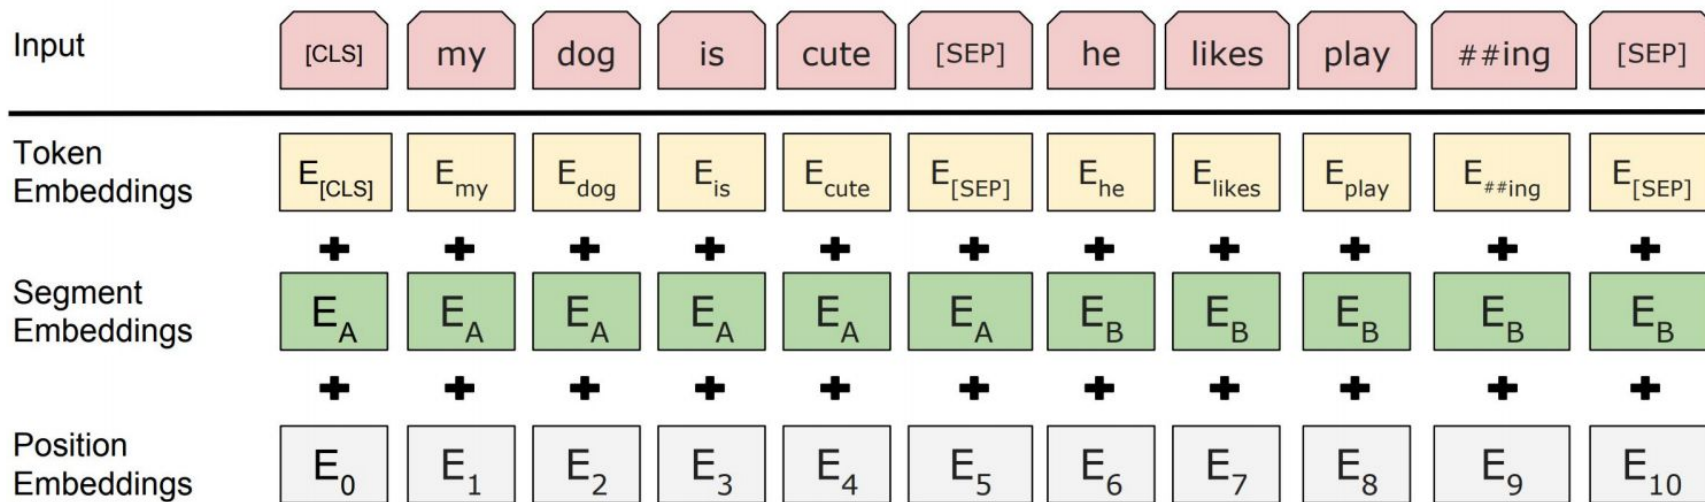


Pre-trained Language Models

- BERT - Bidirectional Encoder Representations from Transformers
- **Main idea:** Propose two pre-training objectives for bidirectional encoders
 - **MASKED LANGUAGE MODEL**
 - **NEXT SENTENCE PREDICTION**
- **Strengths:**
 - Just fine-tune BERT model for downstream tasks to outperform many heavily-engineered tasks
 - BERT advances the state of the art for eleven NLP tasks
 - question answering, textual entailment, sentence similarity

INPUT REPRESENTATION

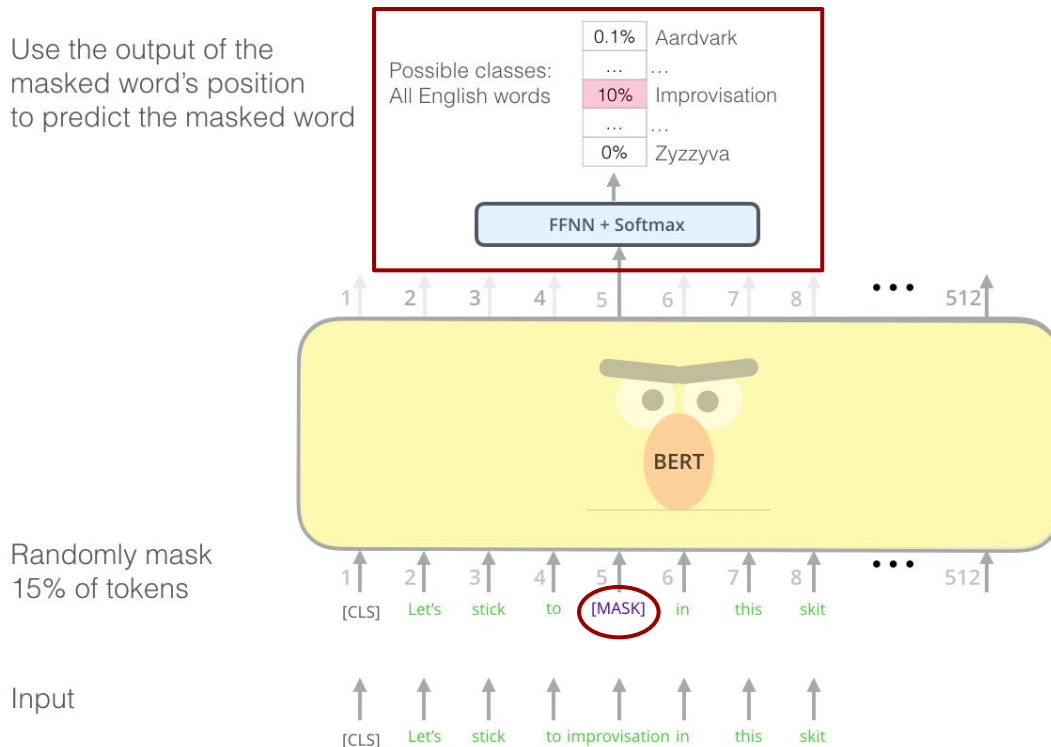
- **WordPiece** Tokenizer w/ [CLS] token, subword representation



MASKED LANGUAGE MODEL

- **Idea:** **Predict** the original word of a **masked word** based only on its context

Use the output of the masked word's position to predict the masked word



MASKING STRATEGY

- 15% of the words are masked at random
 - Tôi muốn đến [cửa hàng] để mua một [thùng] sữa
 - Tôi muốn đến [MASK] để mua một [MASK] sữa
- Too little masking? Too expensive to train
- Too much masking? Not enough context
- Problem with MASK tokens?
 - Mask token never seen at fine-tuning

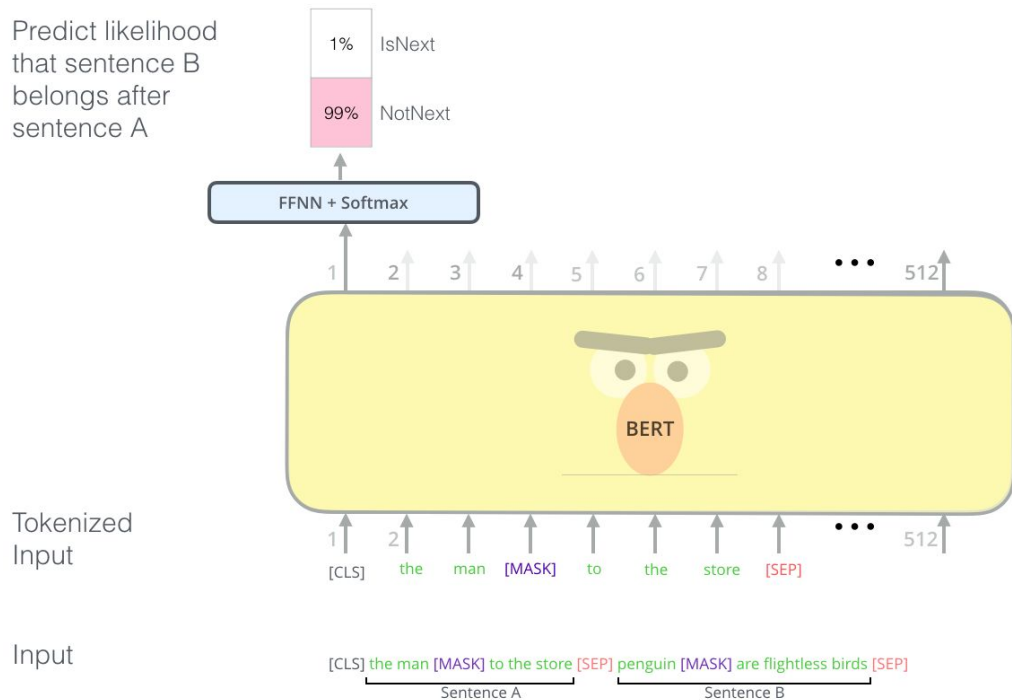
MASKING STRATEGY

- Do not always replace the masked words with **[MASK]** token
- Example: "**Chó** trông rất đáng yêu"
 - **80%** were replaced by the **[MASK] token**: "Chó trông rất đáng **[MASK]**"
 - **10%** were replaced by a **random token**: "Chó trông rất đáng **chuối**"
 - **10%** is **unchanged**: "Chó trông rất đáng **yêu**"

NEXT SENTENCE PREDICTION

- Many downstream tasks are based on understanding the **relationship between two sentences**
 - Question Answer (QA) and Natural Language Inference (NLI)
- LM does not directly capture that relationship
- 50% of training data is “consecutive”

Predict likelihood that sentence B belongs after sentence A



EVALUATION FOR BERT: GLUE

System	MNLI-(m/mm) 392k	QQP 363k	QNLI 108k	SST-2 67k	CoLA 8.5k	STS-B 5.7k	MRPC 3.5k	RTE 2.5k	Average -
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
BERT _{BASE}	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
BERT _{LARGE}	86.7/85.9	72.1	92.7	94.9	60.5	86.5	89.3	70.1	82.1

Table 1: GLUE Test results, scored by the evaluation server (<https://gluebenchmark.com/leaderboard>). The number below each task denotes the number of training examples. The “Average” column is slightly different than the official GLUE score, since we exclude the problematic WNLI set.⁸ BERT and OpenAI GPT are single-model, single task. F1 scores are reported for QQP and MRPC, Spearman correlations are reported for STS-B, and accuracy scores are reported for the other tasks. We exclude entries that use BERT as one of their components.

MODEL DETAILS

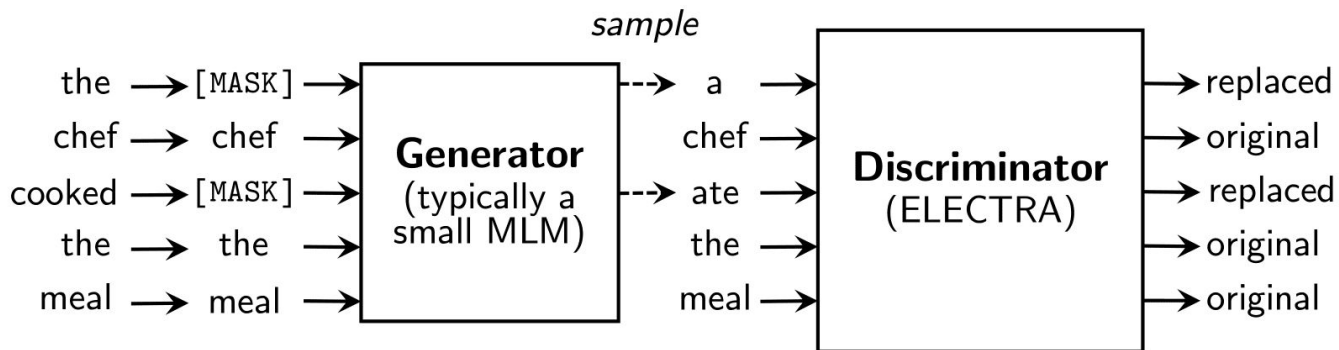
Data	Wikipedia (2.5B words) + BookCorpus (800M words)	
Batch Size	131,072 words (1024 sequences * 128 length or 256 sequences * 512 length)	
Training Time	1M steps (~40 epochs)	
Optimizer	AdamW, 1e-4 learning rate, linear decay	
Architecture	Using Transformer Encoder stack	
	BERT-Base: 12-layer, 768-hidden, 12-head, 110M-params	BERT-Large: 24-layer, 1024-hidden, 16-head, 340M-params

BERT VARIANTS

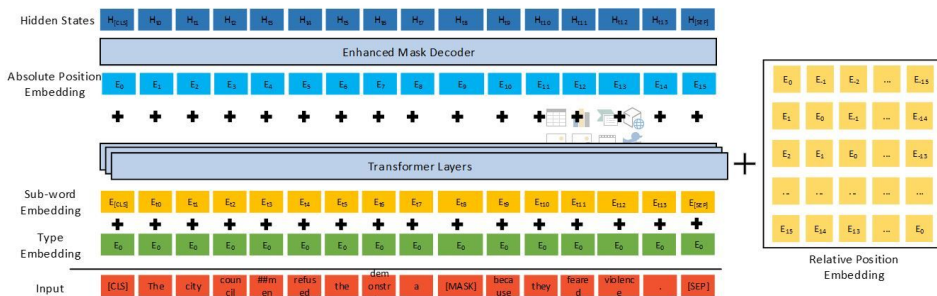
Comparison	BERT	RoBERTa	ELECTRA	DeBERTa
Data	16GB	160GB	16GB for small & base	160GB
Pre-trained Task	Masked Language Model Next Sentence Prediction	Masked Language Model Using Dynamic Masking	Replaced Token Detection	Replaced Token Detection w/ Efficient Embedding Sharing
Parameters	Base: 110M Large: 340M	Base: 125M Large: 355M	Small: 14M Base: 110M Large: 335M	Small: 44M Base: 98M Large: 304M
Tokenization	WordPiece	Byte-Pair Encoding (BPE)	WordPiece	Byte-Pair Encoding (BPE)
Training	Described in slide 20	Larger batchsize, longer sequences input, longer training time	Faster training	500K steps

Efficiently Learning an Encoder that Classifies Token Replacements Accurately


→ Every output position is used




- Disentangled Attention
- Replaced Token Detection with Efficient Embedding Sharing



Standard
Transformer




DeBERTa




LONG TEXT

- Longformer (Beltagy et al. 2020): combine (local) sliding window and global attention
- BigBird (Zaheer et al. 2021): use sparse (global) attention mechanism
- RoFormer (Su et al. 2021): use Rotary Position Embedding




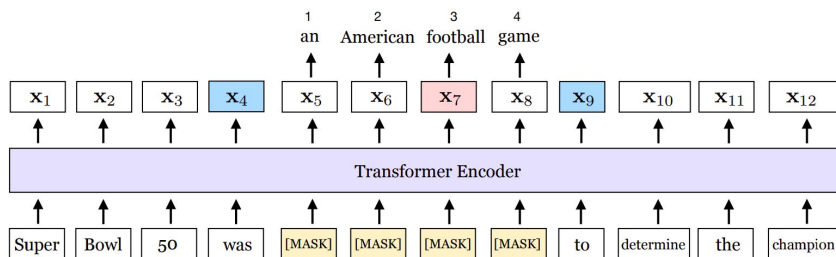
COMPACT PRE-TRAINED MODELS

- DistilBERT: Train a model to match the distribution of BERT (Knowledge Distillation)
- TinyBERT: Same as above but uses layer-wise KD.
- MobileBERT: Same as above but uses inverted-bottleneck and blockwise KD.
- ALBERT: Smaller embeddings, and parameter sharing across all layers.

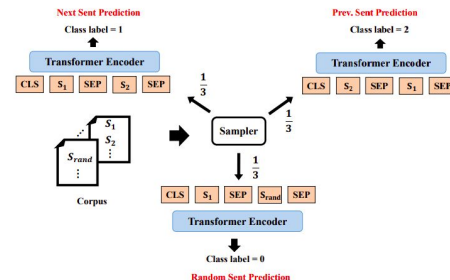


MORE TRAINING OBJECTIVES

- Whole Word Masking
- SpanBERT (Joshi et al., 2019) → Span Masking + Span Boundary Objective
- StructBERT (Wang et al., 2019) → Word-level Objective + Sentence-level Objective



(a) Word Structural Objective



(b) Sentence Structural Objective

Vietnamese BERT

Comparison	PhoBERT	viELECTRA	ViDeBERTa
Data	20GB (Wiki+News)	60GB (Oscar+News)	138GB CC100
Pre-trained Task	Masked Language Model	Replaced Token Detection	Replaced Token Detection w/ Efficient Embedding Sharing
Parameters	Base Large	Base: 125M	Xsmall: 22M Base: 86M Large: 304M
Tokenization	Word-level fastBPE	Word Piece*	Word-level BPE
Training	256len/500K steps	100 epochs	500K steps

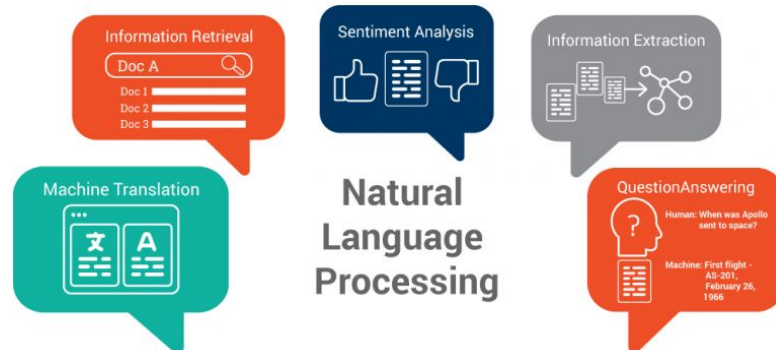
Downstream Tasks

EVALUATION FOR BERT: GLUE

- General Language Understanding Evaluation (**GLUE**) benchmark: Standard split of data to train, validation, test, where labels for the test set is only held in the server.
 - <https://gluebenchmark.com/>
- Sentence pair classification
 - **MNLI**, Multi-Genre Natural Language Inference
 - **QQP**, Quora Question Pairs
 - **QNLI**, Question Natural Language Inference
 - **STS-B**, The Semantic Textual Similarity Benchmark
 - **MRPC**, Microsoft Research Paraphrase Corpus
 - **RTE**, Recognizing Textual Entailment
 - **WNLI**, Winograd NLI is a small natural language inference dataset
- Single sentence classification
 - **SST-2**, The Stanford Sentiment Treebank
 - **CoLA**, The Corpus of Linguistic Acceptability

COMMON DOWNSTREAM TASKS

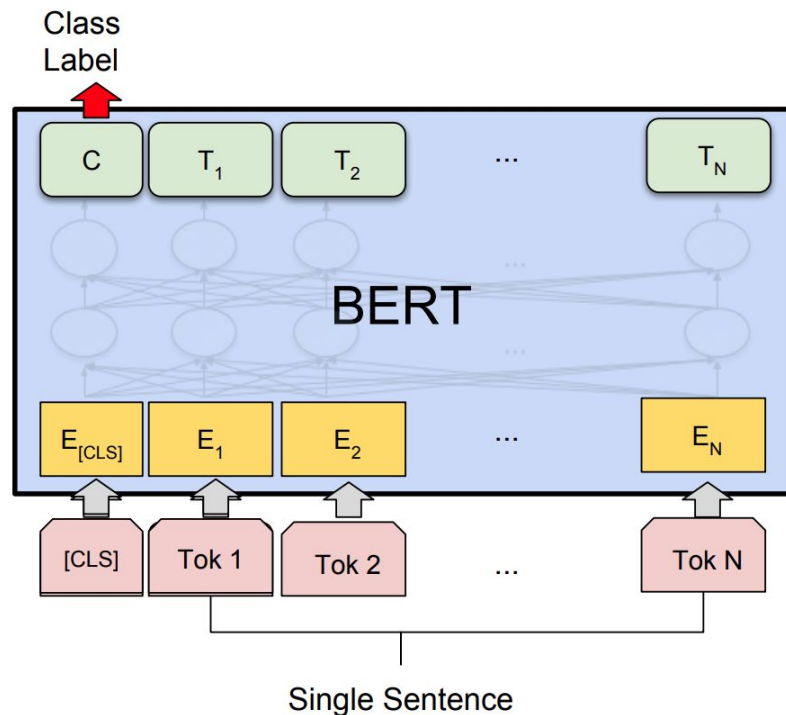
- Sentiment Analysis
- Named Entity Recognition
- Natural Language Inference
- Question Answering



SINGLE SENTENCE CLASSIFICATION

- **Sentiment Analysis**
 - Tôi thích đọc sách -> Positive
 - Tôi ghét phim kinh dị -> Negative
- **Input:** [CLS] Tôi thích đọc sách [SEP] -> Positive/Negative?

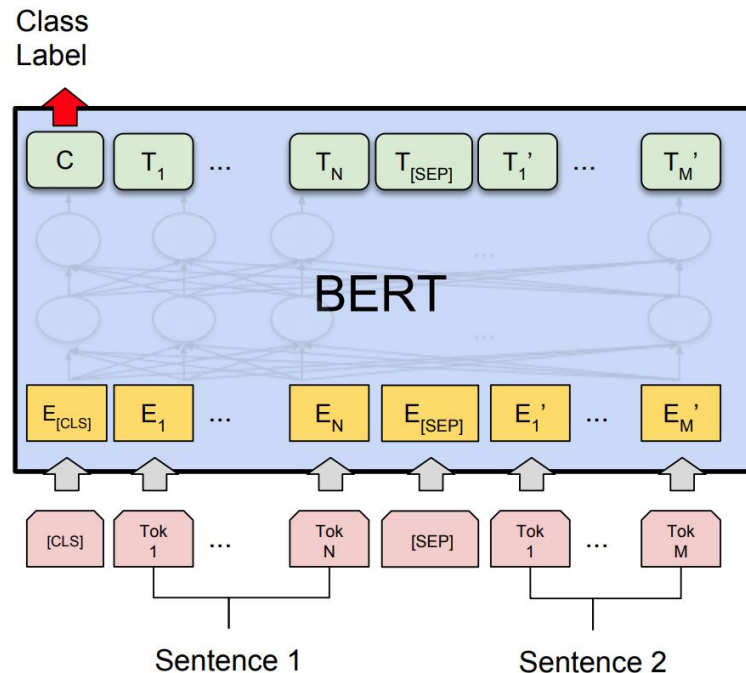
Other applications: Aspect-based Sentiment Analysis, Spam Detection, etc.



SENTENCE-PAIR CLASSIFICATION

- **Natural Language Inference**
 - Sent 1: Tôi làm nhiều việc tốt
 - Sent 2: Tôi được mọi người yêu quý
 - Sent 1 -> Sent 2: Entailment
- **Input:** [CLS] Tôi làm nhiều việc tốt [SEP] Tôi được mọi người yêu quý [SEP] -> Entailment/Contradiction/Neutral?

Other applications: Semantic Similarity, Ranking, etc.

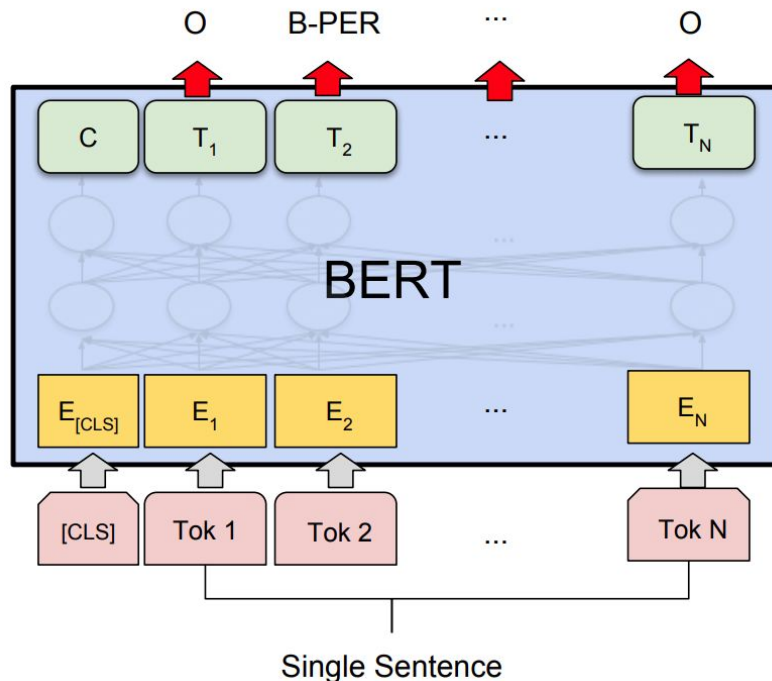


TOKEN CLASSIFICATION

- **Named Entity Recognition (NER)**: detect entities such as PER, LOC, ORG
- **Text**: Mark Zuckerberg cho ra đời Facebook trong phòng ký túc xá của mình tại Đại học Harvard vào ngày 4 tháng 2 năm 2004.
- **Input**: [CLS] Text [SEP]
- **Output**: B-PER, I-PER, O, O,...

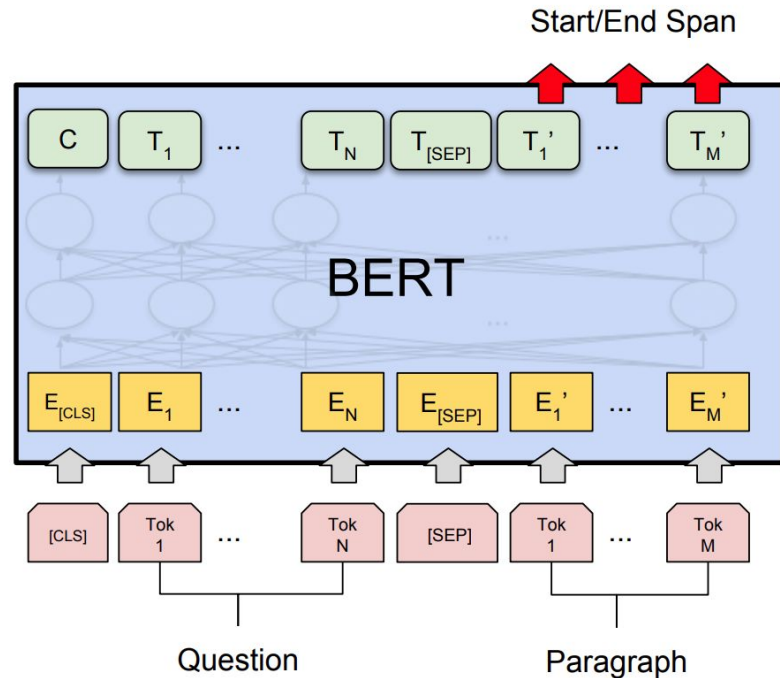
Other applications: POS Tagging, Punctuation Prediction, etc.

Source: <https://arxiv.org/pdf/1810.04805.pdf>

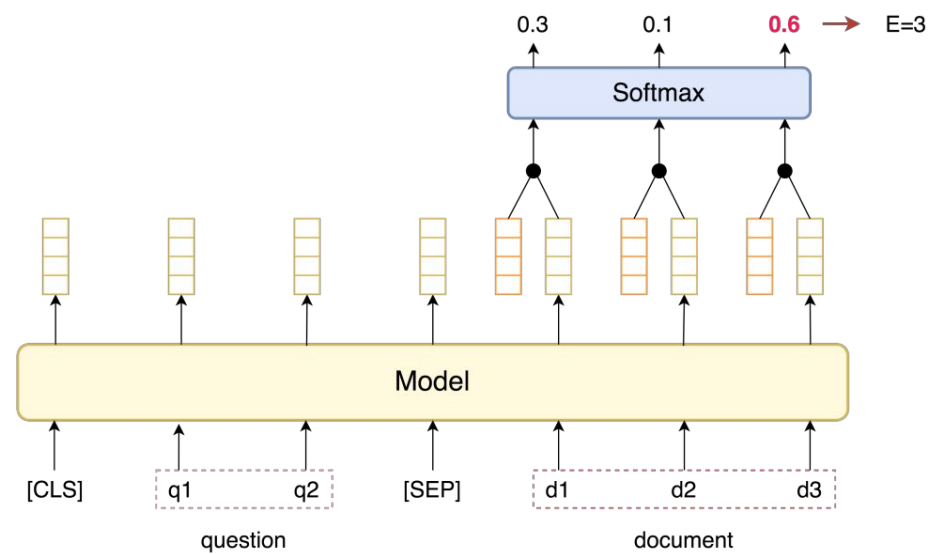
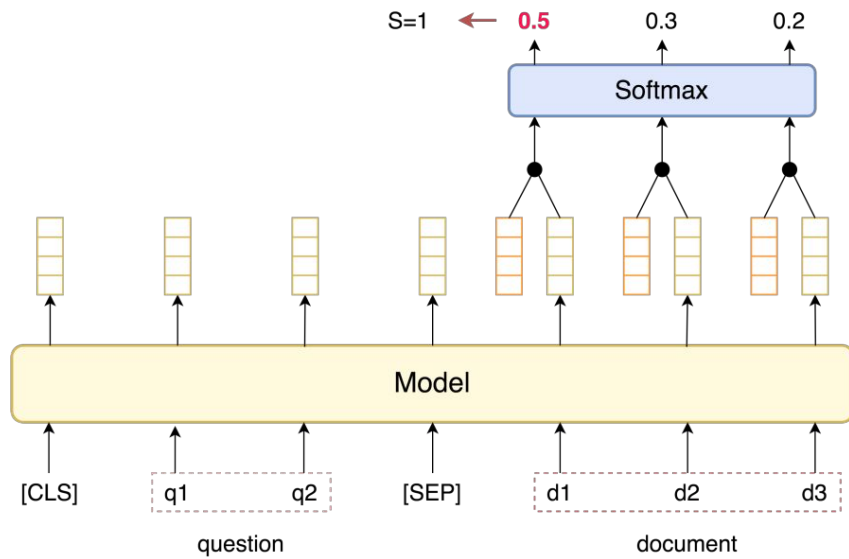


QUESTION ANSWERING

- **Detect answer span** in paragraph
- **Question:** Đây là thủ đô của Việt Nam?
- **Paragraph:** Hà Nội là thủ đô của nước Cộng hoà Xã hội chủ nghĩa Việt Nam, cũng là kinh đô của...
- **Input:** [CLS] Question [SEP] Paragraph [SEP]
- **Output:** start=0, end=2
- **Answer:** Hà Nội

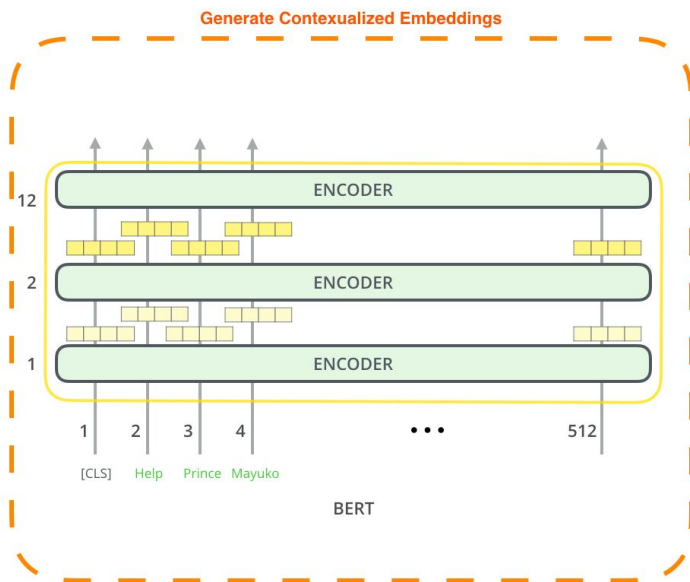


QUESTION ANSWERING




FEATURE EXTRACTION

- Which layer should be used?



The output of each encoder layer along each token's path can be used as a feature representing that token.












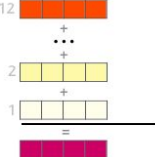

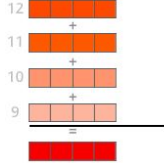



But which one should we use?

FEATURE EXTRACTION

- Depend on the task



What is the best contextualized embedding for “Help” in that context?
For named-entity recognition task CoNLL-2003 NER

		Dev F1 Score
12		
...		
7		
6		
5		
4		
3		
2		
1		
		
	Help	
First Layer	Embedding 	91.0
Last Hidden Layer	12 	94.9
Sum All 12 Layers		95.5
Second-to-Last Hidden Layer	11 	95.6
Sum Last Four Hidden		95.9
Concat Last Four Hidden		96.1

Source: <https://jalammar.github.io/illustrated-bert/>

FINE-TUNING STRATEGY

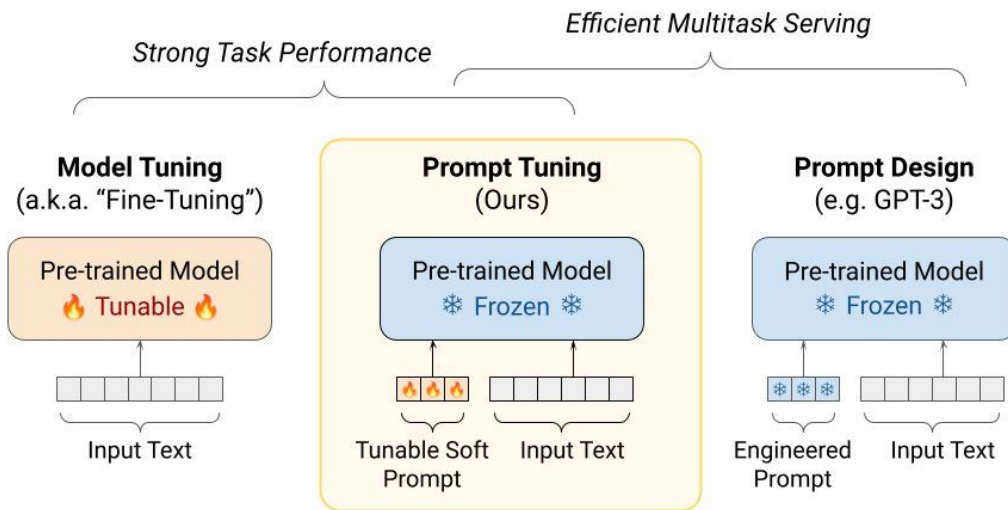

- Train the entire architecture
- Train some layers while freezing others
- Freeze all layers
- Freeze all - train an adapter



Why doesn't freeze all BERT layers during fine-tuning tasks?

PROMPT TUNING

Instead of fine-tuning, hold the model fixed and tune a fixed-size vector representation.



WHAT WE LEARN TODAY

- The different between word embeddings: Word2Vec, GloVe, CoVe, ELMo
- Types of learning: multi-task, transfer learning, pretraining.
- The architecture of BERT and its variants
- Introduction of downstream tasks: NLI, QA, NER, etc

Thank you!
